

Face mask detection

Contents

- [Step 0](#): Import libraries
- [Step 1](#): Load and explore data
- [Step 2](#): Load data using Keras function
- [Step 3](#): Analyze training data
- [Step 4](#): Build the model
- [Step 5](#): Fit the model
- [Step 6](#): Evaluate the model
- [Step 7](#): Results
- [Step 8](#): Improve the model (optional)
- [Extra](#): Test own photos

Step 0: Import libraries

```
In [2]: import PIL.Image
import matplotlib.pyplot as plt
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential

import kagglehub
import pathlib

import numpy as np
```

Step 1: Load and explore data

```
In [3]: # Download latest version
#path = kagglehub.dataset_download("ashishjangra27/face-mask-12k-images-dataset")

#print("Path to dataset files:", path)
```

```
In [4]: path = 'C:/Users/ronal/.cache/kagglehub/datasets/ashishjangra27/face-mask-12k-images-dataset'
```

```
In [5]: path = pathlib.Path(path).with_suffix('')./'Face Mask Dataset'
```

```
In [7]: len(list(path.glob("*/*/*.png")))
```

```
Out[7]: 11792
```

```
In [8]: mask = list(path.glob("Train/WithMask/*"))  
PIL.Image.open(str(mask[0]))
```

Out[8]:



```
In [9]: PIL.Image.open(str(mask[100]))
```

Out[9]:



```
In [10]: no_mask = list(path.glob("Train/WithoutMask/*"))  
PIL.Image.open(str(no_mask[0]))
```

Out[10]:



```
In [11]: PIL.Image.open(str(no_mask[100]))
```

Out[11]:



Step 2: Load data using Keras function

```
In [12]: batch_size = 32  
img_height = 128  
img_width = 128
```

```
In [13]: train_ds = tf.keras.utils.image_dataset_from_directory(  
    path/'Train',  
    seed=9,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```

Found 10000 files belonging to 2 classes.

```
In [14]: val_ds = tf.keras.utils.image_dataset_from_directory(  
    path/'Validation',  
    seed=9,
```

```
image_size=(img_height, img_width),  
batch_size=batch_size)
```

Found 800 files belonging to 2 classes.

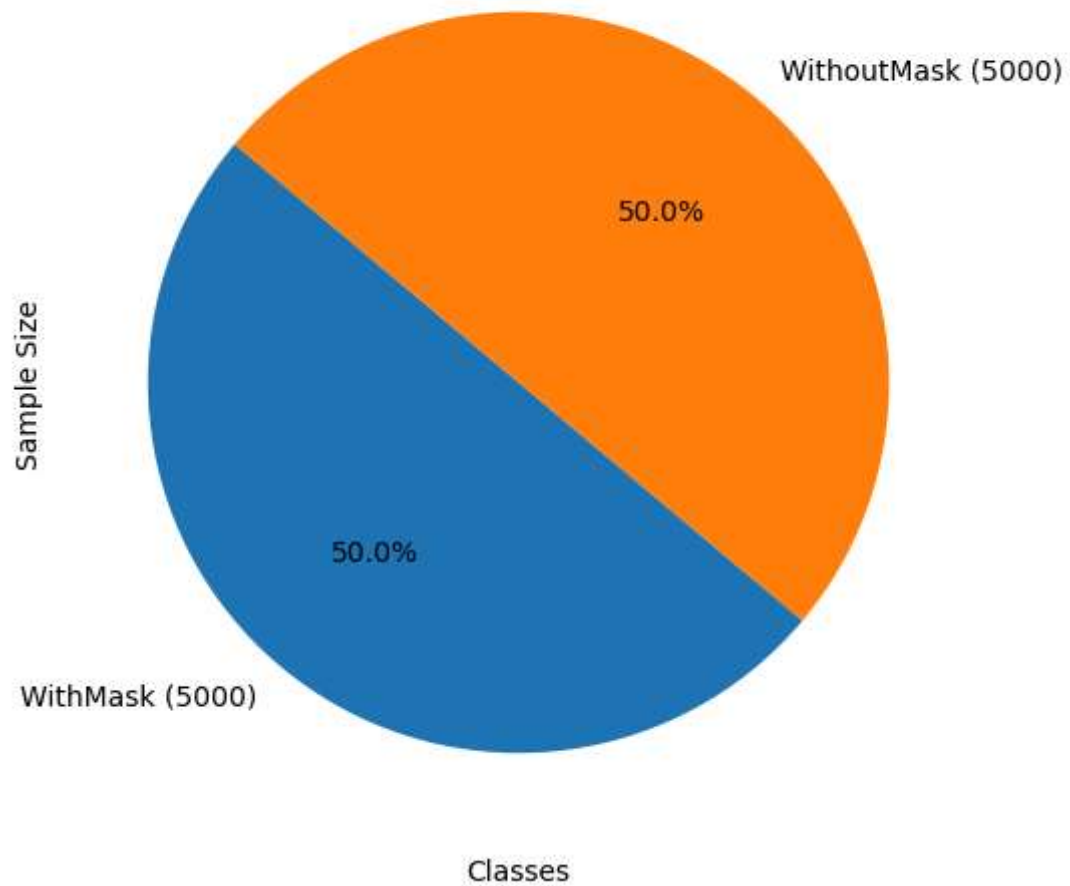
```
In [15]: test_ds = tf.keras.utils.image_dataset_from_directory(  
        path/'Test',  
        seed=9,  
        batch_size=32,  
        image_size=(img_height, img_width))
```

Found 992 files belonging to 2 classes.

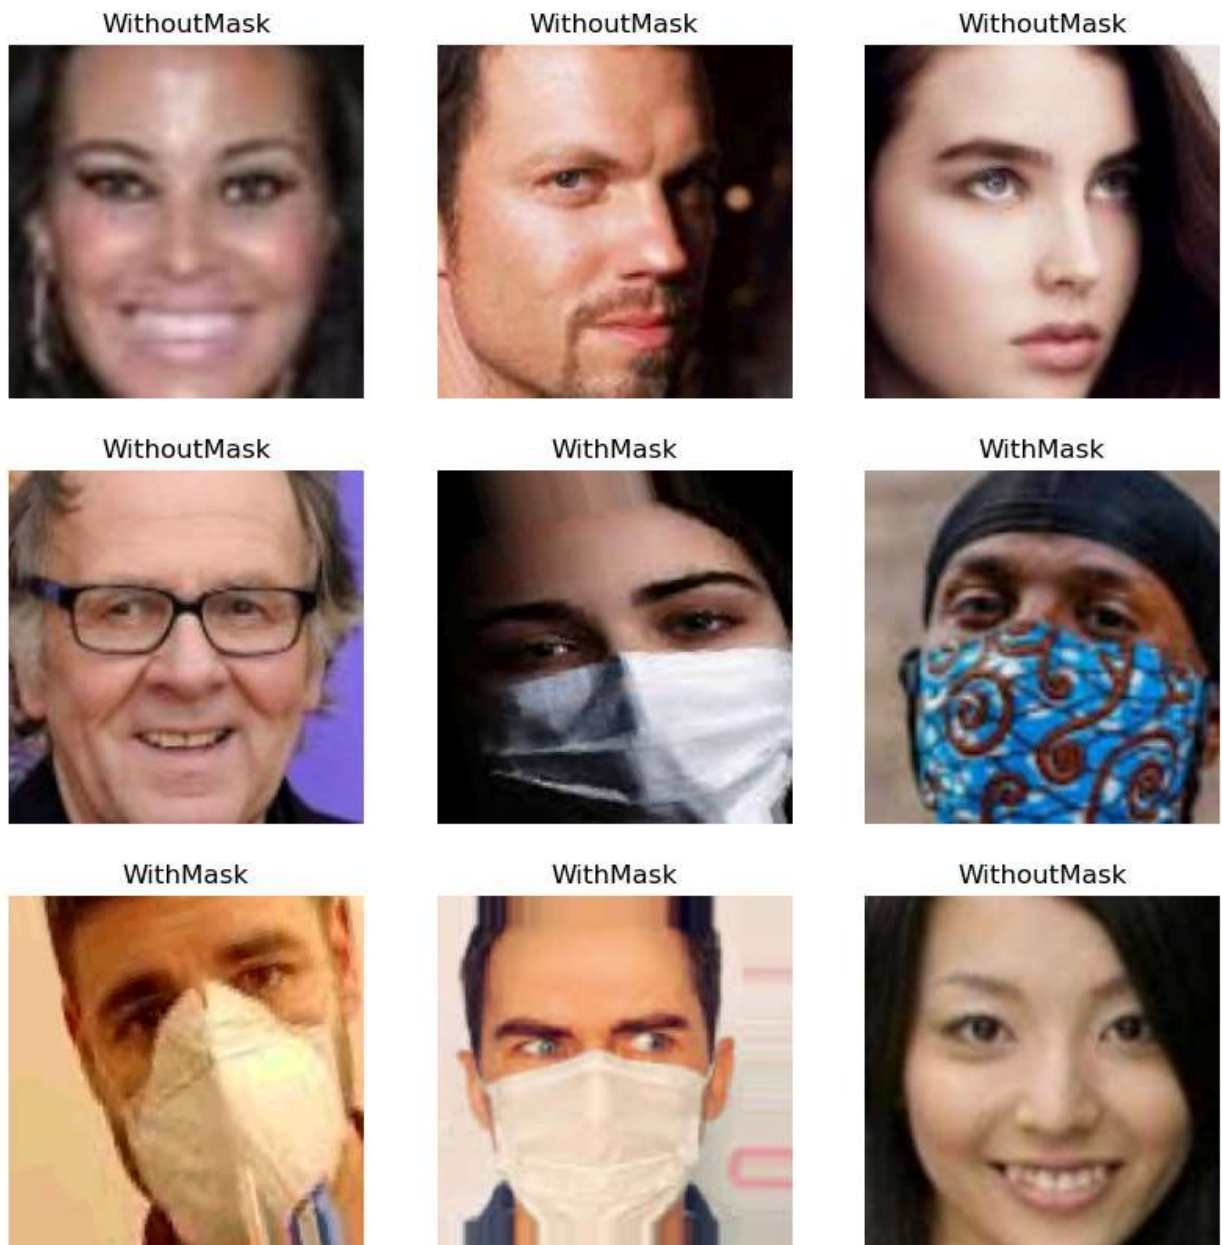
Step 3: Analyze training data

```
In [16]: class_labels = train_ds.class_names  
        class_counts = [len(mask), len(no_mask)]  
  
        labels = [f"{name} ({count})" for name, count in zip(class_labels, class_counts)]  
  
        plt.figure(figsize=(8, 6))  
        plt.pie(class_counts, labels=labels, autopct='%1.1f%%', startangle=140)  
        plt.xlabel('Classes')  
        plt.ylabel('Sample Size')  
        plt.title('Distribution of training sample size')  
        plt.show()
```

Distribution of training sample size



```
In [17]: plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_labels[labels[i]])
        plt.axis("off")
```



Step 4: Build the model

```
In [18]: AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
In [19]: num_classes = len(class_labels)

model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
```

```

layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(num_classes)
])

```

c:\Users\ronal\anaconda3\Lib\site-packages\keras\src\layers\preprocessing\tf_data_layer.py:19: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

```

In [20]: model.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])

```

```
In [21]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 128, 128, 3)	0
conv2d (Conv2D)	(None, 128, 128, 16)	448
max_pooling2d (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_1 (Conv2D)	(None, 64, 64, 32)	4,640
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_2 (Conv2D)	(None, 32, 32, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 64)	0
flatten (Flatten)	(None, 16384)	0
dense (Dense)	(None, 128)	2,097,280
dense_1 (Dense)	(None, 2)	258

Total params: 2,121,122 (8.09 MB)

Trainable params: 2,121,122 (8.09 MB)

Non-trainable params: 0 (0.00 B)

```

In [22]: from tensorflow.keras.utils import plot_model

plot_model(model, to_file='model_architecture.png', show_shapes=True, show_layer_names

```

You must install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model` to work.

Step 5: Fit the model

```
In [23]: epochs=10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
```

```
Epoch 1/10
313/313 ————— 46s 121ms/step - accuracy: 0.8780 - loss: 0.2711 - val_a
ccuracy: 0.9875 - val_loss: 0.0470
Epoch 2/10
313/313 ————— 35s 112ms/step - accuracy: 0.9847 - loss: 0.0385 - val_a
ccuracy: 0.9875 - val_loss: 0.0324
Epoch 3/10
313/313 ————— 36s 115ms/step - accuracy: 0.9873 - loss: 0.0351 - val_a
ccuracy: 0.9925 - val_loss: 0.0255
Epoch 4/10
313/313 ————— 34s 108ms/step - accuracy: 0.9915 - loss: 0.0230 - val_a
ccuracy: 0.9950 - val_loss: 0.0211
Epoch 5/10
313/313 ————— 33s 106ms/step - accuracy: 0.9930 - loss: 0.0170 - val_a
ccuracy: 0.9912 - val_loss: 0.0204
Epoch 6/10
313/313 ————— 33s 104ms/step - accuracy: 0.9921 - loss: 0.0178 - val_a
ccuracy: 0.9912 - val_loss: 0.0219
Epoch 7/10
313/313 ————— 38s 122ms/step - accuracy: 0.9964 - loss: 0.0109 - val_a
ccuracy: 0.9950 - val_loss: 0.0246
Epoch 8/10
313/313 ————— 38s 121ms/step - accuracy: 0.9972 - loss: 0.0084 - val_a
ccuracy: 0.9887 - val_loss: 0.0379
Epoch 9/10
313/313 ————— 38s 121ms/step - accuracy: 0.9924 - loss: 0.0244 - val_a
ccuracy: 0.9912 - val_loss: 0.0183
Epoch 10/10
313/313 ————— 40s 128ms/step - accuracy: 0.9970 - loss: 0.0077 - val_a
ccuracy: 0.9900 - val_loss: 0.0282
```

```
In [24]: model.save('facemask_model_simple.keras')
```

```
In [25]: #from tensorflow.keras.models import load_model

#model = load_model('facemask_model_simple.keras')
#predictions = model.predict(test_ds)
```

Step 6: Evaluate the model

```
In [26]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

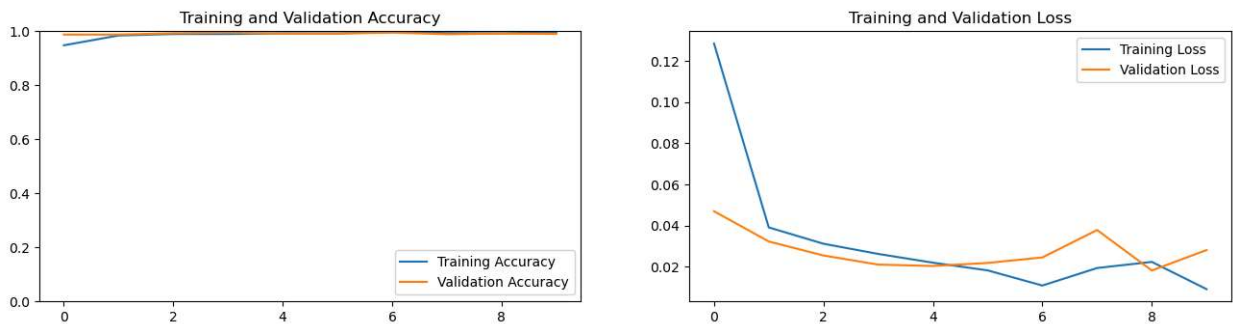
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(10)

plt.figure(figsize=(16, 8))
```

```
plt.subplot(2, 2, 1)
plt.ylim(0,1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [36]: # Evaluate the model on the test dataset
test_loss, test_accuracy = model.evaluate(test_ds)

# Print the test Loss and accuracy
print("Test Loss:", test_loss)
print("Test Accuracy:", test_accuracy)
```

```
31/31 ————— 1s 38ms/step - accuracy: 0.9858 - loss: 0.0492
Test Loss: 0.051358141005039215
Test Accuracy: 0.9868951439857483
```

Step 7: Results

```
In [28]: predictions = model.predict(test_ds)
```

```
31/31 ————— 2s 44ms/step
```

```
In [37]: plt.figure(figsize=(10, 10))

num_images_to_plot = 25 # Limit of images to display
current_plot = 0 # Counter for images displayed

for images, labels in test_ds.take(19):
    if current_plot >= num_images_to_plot:
        break # Stop when the plot limit is reached

    # Get predictions for the current batch
    batch_predictions = model(images, training=False)

    # Loop over each image in the batch
    for image, prediction in zip(images, batch_predictions):
        if current_plot >= num_images_to_plot:
            break # Stop if we have reached the plot limit
```



```

# Prepare image and prediction details
image = image.numpy().astype("uint8")
prediction_score = tf.nn.softmax(prediction)
predicted_label = class_labels[tf.argmax(prediction_score).numpy()]
confidence = tf.reduce_max(prediction_score).numpy()

# Plot image with label and confidence
plt.subplot(5, 5, current_plot + 1)
plt.imshow(image)
plt.axis("off")
plt.title(f"{predicted_label} ({confidence:.2f})")

current_plot += 1 # Increment plot count

plt.tight_layout()
plt.show()

```

WithoutMask (0.99)



WithoutMask (1.00)



WithoutMask (1.00)



WithMask (0.96)



WithMask (1.00)



WithoutMask (1.00)



WithMask (1.00)



WithMask (1.00)




WithoutMask (0.98)



WithoutMask (1.00)




WithoutMask (0.99)



WithMask (1.00)



WithoutMask (0.99)



WithoutMask (0.99)



WithoutMask (0.99)



WithMask (1.00)



WithMask (1.00)



WithMask (0.99)



WithoutMask (1.00)



WithoutMask (0.99)



WithMask (1.00)



WithoutMask (1.00)



WithMask (1.00)



WithMask (1.00)



WithMask (1.00)



Step 8: Improve model (optional)

```
In [30]: data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal",  
                           input_shape=(img_height,  
                                         img_width,  
                                         3)),  
        layers.RandomRotation(0.1),  
        layers.RandomZoom(0.1),  
    ]  
)
```

```
In [ ]: plt.figure(figsize=(10, 10))  
for images, _ in train_ds.take(1):  
    for i in range(9):  
        augmented_images = data_augmentation(images)  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(augmented_images[0].numpy().astype("uint8"))  
        plt.axis("off")
```

```
In [32]: model = Sequential([  
    data_augmentation,  
    layers.Rescaling(1./255),  
    layers.Conv2D(16, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(32, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(64, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Dropout(0.2),  
    layers.Flatten(),  
    layers.Dense(128, activation='relu'),  
    layers.Dense(num_classes, name="outputs")  
)
```

```
In [33]: model.compile(optimizer='adam',  
                       loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
                       metrics=['accuracy'])
```

```
In [ ]: model.summary()
```

```
In [ ]: epochs = 10  
history = model.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=epochs  
)
```

```
In [ ]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']  
  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
  
epochs_range = range(epochs)
```

```

plt.figure(figsize=(16, 8))
plt.subplot(2, 2, 1)
plt.ylim(0,1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```

```

In [ ]: plt.figure(figsize=(10, 10))

# Keep track of how many images are plotted
num_images_to_plot = 25
current_plot = 0

for batch_index, (images, labels) in enumerate(test_ds):
    # Perform predictions on the current batch of images
    batch_predictions = model(images, training=False) # Get predictions directly from

    for image_index in range(images.shape[0]):
        if current_plot >= num_images_to_plot:
            break # Stop if we have plotted 25 images

        plt.subplot(5, 5, current_plot + 1)

        # Extract the image
        image = images[image_index].numpy().astype("uint8")

        # Get the prediction and score for the current image
        prediction_score = tf.nn.softmax(batch_predictions[image_index])
        predicted_label = class_labels[np.argmax(prediction_score)]
        confidence = tf.reduce_max(prediction_score).numpy() # Extract the highest co

        # Plot the image
        plt.imshow(image)
        plt.axis("off")

        # Set the title with the predicted label and confidence
        plt.title(f"{predicted_label} ({confidence:.2f})")

        current_plot += 1

    if current_plot >= num_images_to_plot:
        break # Stop after plotting 25 images

plt.tight_layout()
plt.show()

```

Additional: Test own photos

```
In [45]: import cv2
import matplotlib.pyplot as plt

img_path = pathlib.Path('C:/Users/ronal/Downloads/ronald.jpg') #Change path

# Cargar el modelo Haar Cascade para detección de rostros
face_cascade = cv2.CascadeClassifier(cv2.data.harcascades + 'haarcascade_frontalface_

# Cargar la imagen
image = cv2.imread(img_path)
# Convertir la imagen a escala de grises, ya que Haar Cascade funciona mejor en grises
gray_image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Detectar rostros en la imagen
faces = face_cascade.detectMultiScale(gray_image, scaleFactor=1.1, minNeighbors=5, mir

# Dibujar un rectángulo alrededor de cada rostro detectado
for (x, y, w, h) in faces:
    cv2.rectangle(image, (x, y), (x+w, y+h), (255, 0, 0), 2) # Rectángulo azul alrede

# Mostrar la imagen con los rostros detectados
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
plt.axis('off') # Quita los ejes
plt.show()
```



```
In [ ]: # Check if any faces were detected
if len(faces) > 0:
    for (x, y, w, h) in faces:
        # Crop the face portion from the original image
        face = cv2.cvtColor(image[y:y+h, x:x+w], cv2.COLOR_BGR2RGB)
        face = cv2.resize(face, (128, 128))
        # Display the cropped face
        plt.imshow(face)
        plt.axis('off') # Hide axes
```

```
plt.show()

else:
    print("No faces detected.")
```



```
In [ ]: index = 0 #Change index
        face = cv2.cvtColor(image[faces[index][1]:faces[index][1]+faces[index][3], faces[index][2]:faces[index][2]+faces[index][4]], cv2.COLOR_BGR2RGB)
        face = cv2.resize(face, (128, 128))
        plt.imshow(face)
        plt.axis('off')
        plt.show()
```

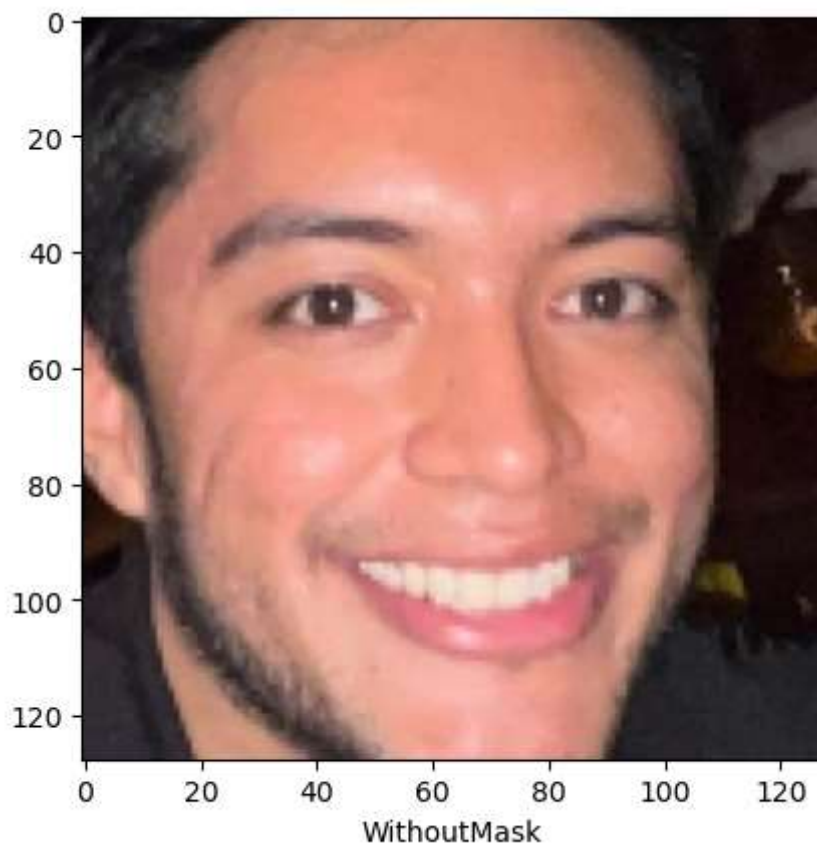


```
In [42]: img_array = tf.expand_dims(face, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

plt.figure()
plt.imshow(face)
plt.grid(False)
plt.xlabel(class_labels[np.argmax(score)])
plt.show()
print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_labels[np.argmax(score)], 100 * np.max(score))
)
```

1/1 ————— 0s 146ms/step



This image most likely belongs to WithoutMask with a 99.63 percent confidence.